



OS Structures

These slides were compiled from the OSC textbook slides (Silberschatz, Galvin, and Gagne) and the instructor's class materials.





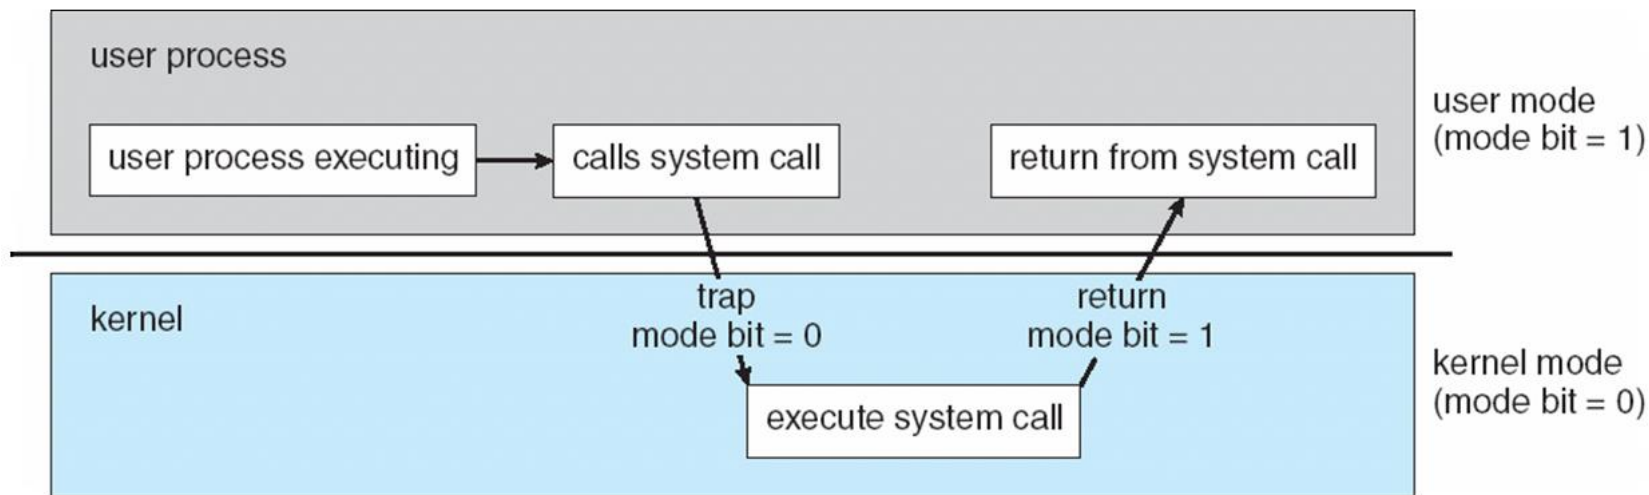
OS Dual-Mode Operations

Provide hardware support to differentiate between at least two modes of operations.

1. **User mode** – execution done on behalf of a user.
2. **Monitor mode** (also *supervisor mode*, *system mode*, or *Kernel mode*) – execution done on behalf of operating system.

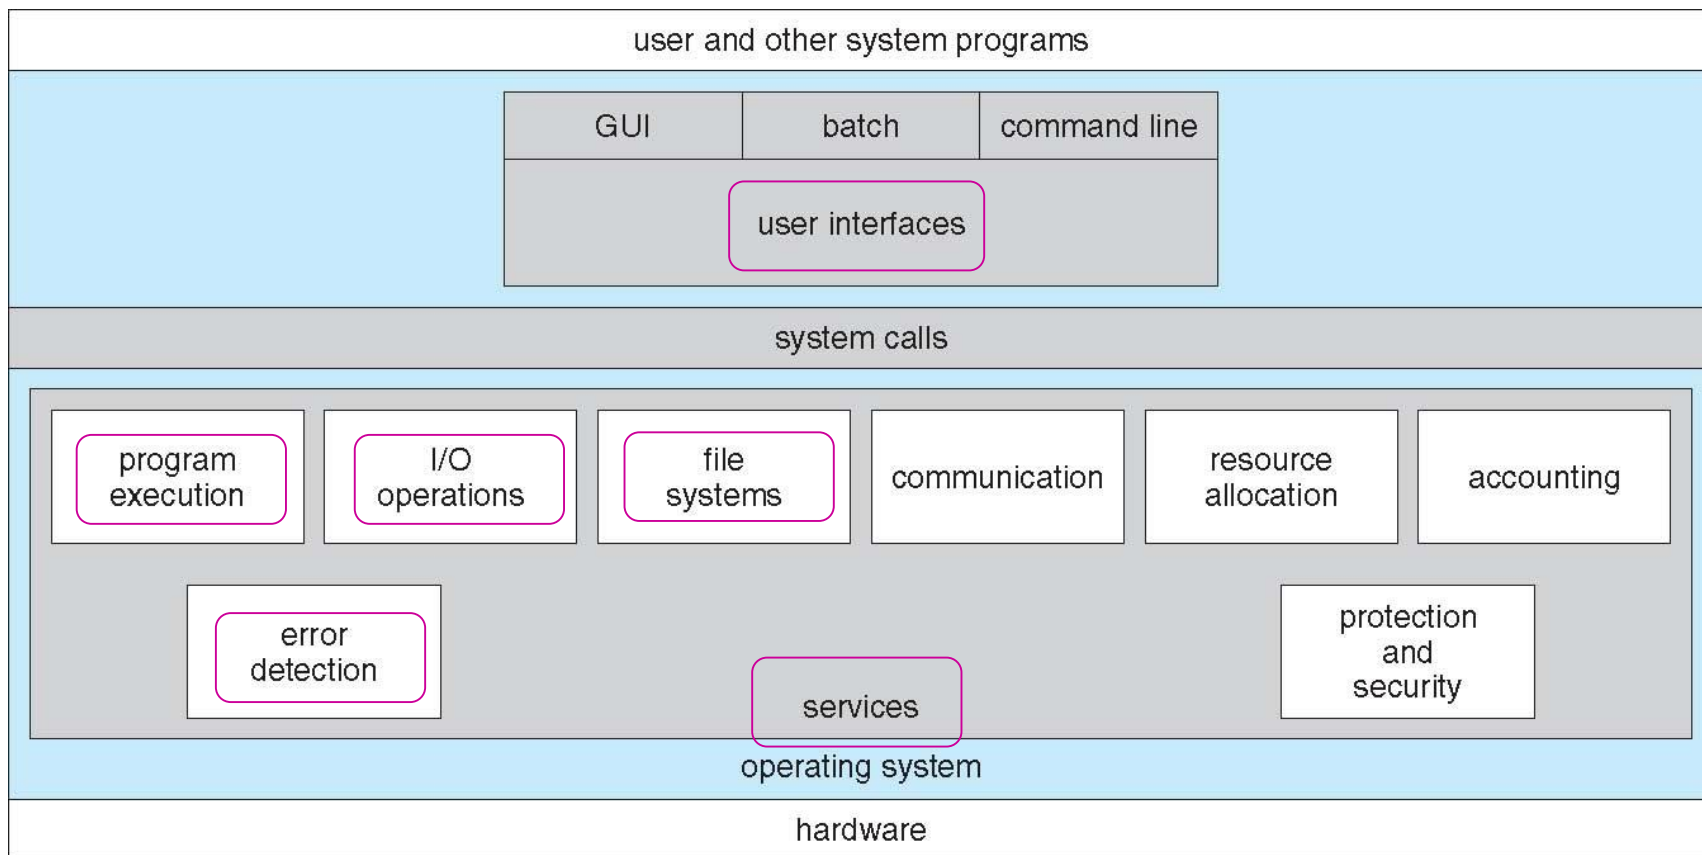
Switching between two modes: transition from **user** -> **kernel** mode

- Device interrupts, hardware traps, system calls cause a trap to the kernel mode
- The operating system returns to the user mode after servicing requests.





A view of OS Services





System Calls and OS “Managements”

- Big picture
 - Process
 - Memory
 - File
 - Others





System Calls

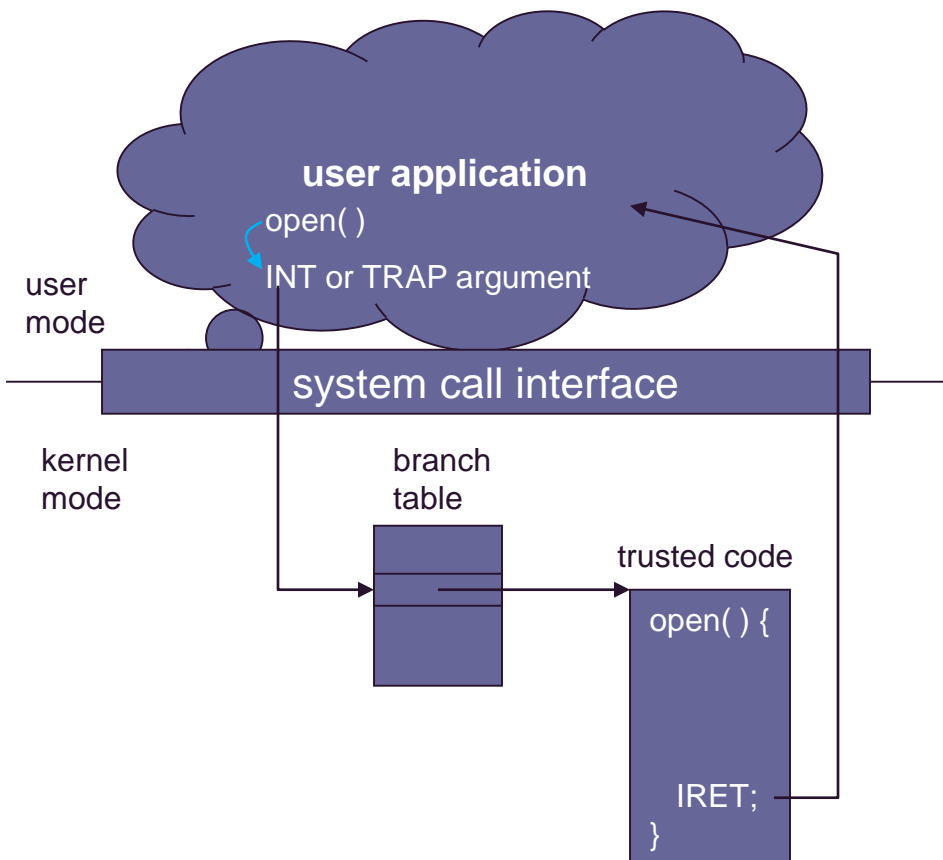
- Programming **interface** to the services provided by the OS
- There is a number associated with each system call (used for indexing)
- Typically written in a high-level language (C or C++), accessed via high-level **API** rather than direct system call use
 - Implementation details hidden
 - Three most common APIs are Win32 API for Windows, POSIX API for POSIX-based systems (for all versions of UNIX), and Java API for the Java virtual machine (JVM)
- Why use APIs rather than system calls?

(Note that the system-call names used throughout this text are generic)





System Calls



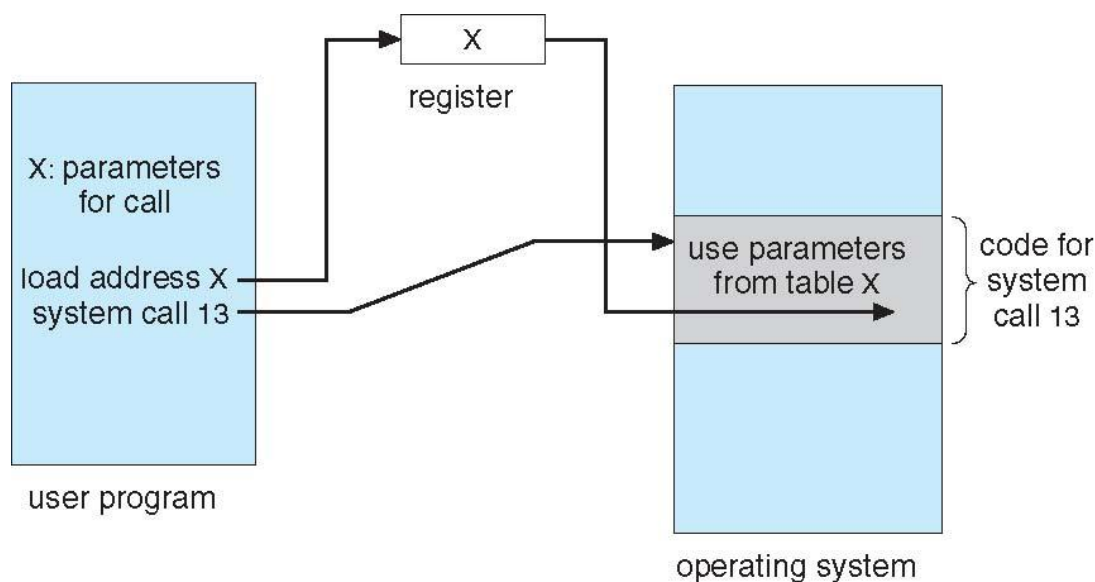
■ Software interface to OS

- A System function is compiled into “trap args”
- “Trap” changes CPU mode from user to kernel.
- CPU checks the OS branch table.
- Jumps to the code pointed to by the table entry.
- Work on the requested call.
- IRET switches from kernel to user





System Calls (Cont'd)

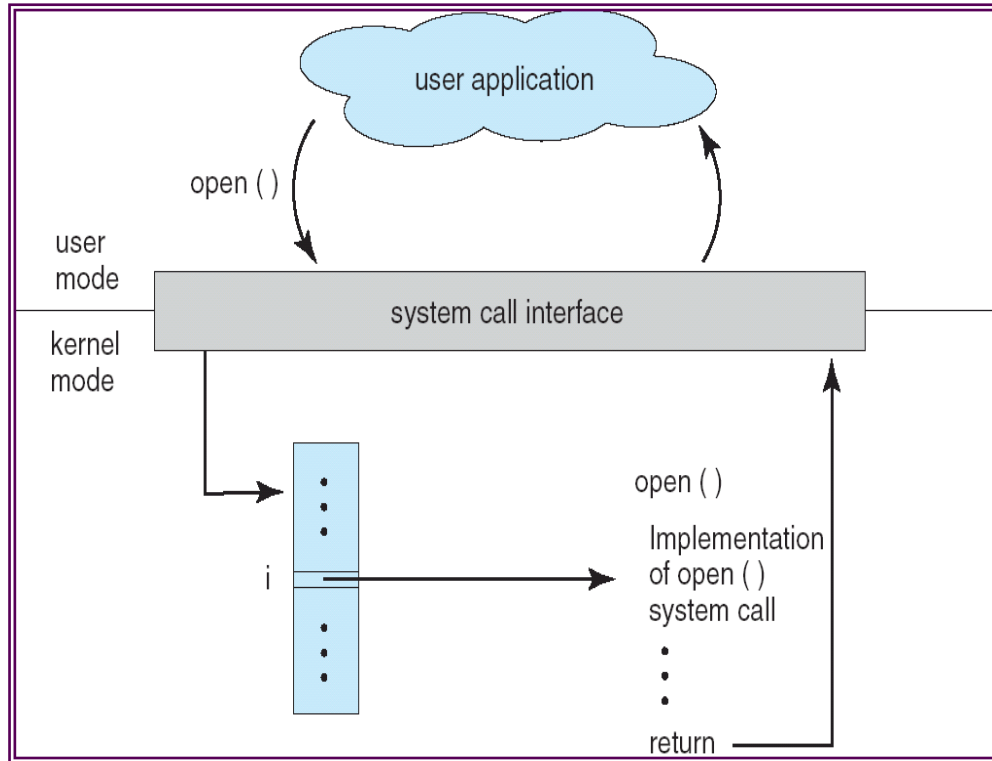


- When a user program executes a **special instruction** like trap
 - CPU recognizes it as a (software) **interrupt**.
 - The mode turns in **kernel** mode.
 - Control jumps to a given **vector** (e.g. 13)
 - The OS **saves** the user program **status**.
 - It then begins to handle the system call.
 - The OS resumes the **registers**.
 - It finally **returns** back to a user program

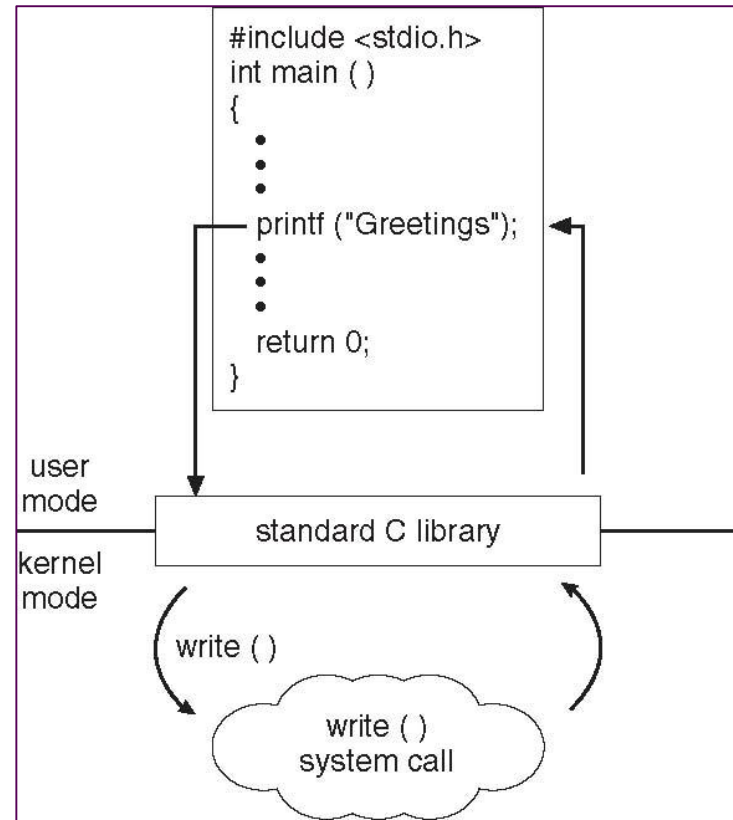




System Calls



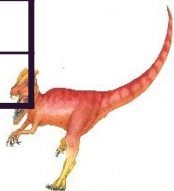
Example: C program invoking `printf()` library call, which calls `write()` system call





System Calls (Cont'd)

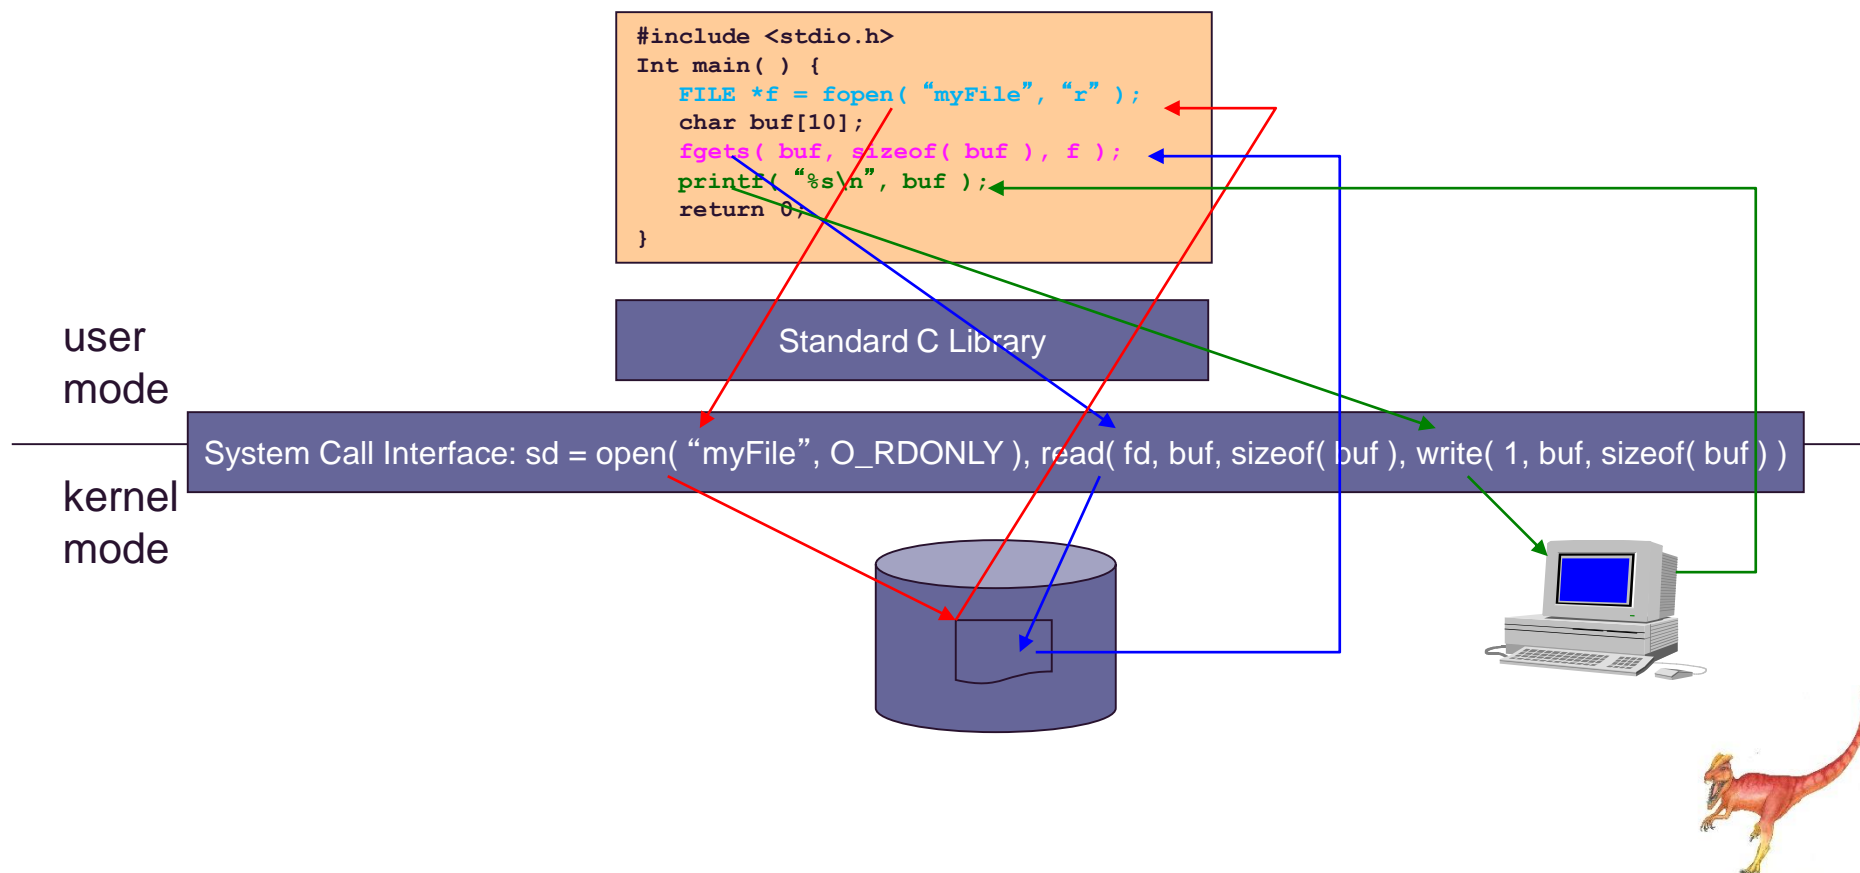
	Windows	Unix
Process Control	CreateProcess()	fork()
	ExitProcess()	exit()
	WaitForSingleObject()	wait()
File Manipulation	CreateFile()	open()
	ReadFile()	read()
	WriteFile()	write()
	CloseHandle()	close()
Device Manipulation	SetConsoleMode()	ioctl()
	ReadConsole()	read()
	WriteConsole()	write()
Information Maintenance	GetCurrentProcessID()	getpid()
	SetTimer()	alarm()
	Sleep()	sleep()
Communication	CreatePipe()	pipe()
	CreateFileMapping()	shmget()
	MapViewOfFile()	mmap()
Protection	SetFileSecurity()	chmod()
	InitializeSecurityDescriptor()	umask()
	SetSecurityDescriptorGroup()	chown()





System Calls Examples

- C program invoking `fopen()`, `fgets()`, and `printf()`, which call `open()`, `read()` and `write()` system calls





Command Interpreters

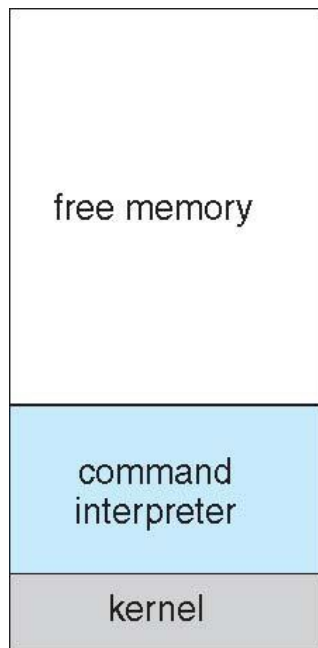
- The program that reads and interprets **control statements**
 - command-line interpreter (in DOS)
 - shell (in UNIX)
 - Mouse-based window and menu system (Windows, Linux, MacOS)

- What control statements can you pass the command interpreter?
 - Program **execution**: a.out, g++, emacs
 - Process **management**: ps, kill, sleep, top, nice, pstack
 - **I/O** operations: lpr, clear, lprm, mt
 - **File**-system manipulation: ls, mkdir, mv, rm, chmod, [u]mount
 - **Communication**: write, ping, mesg

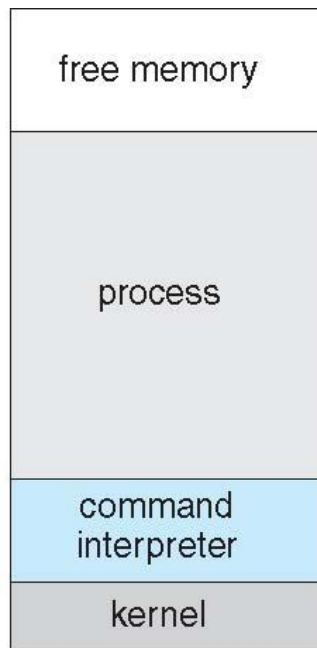




Command Interpreters



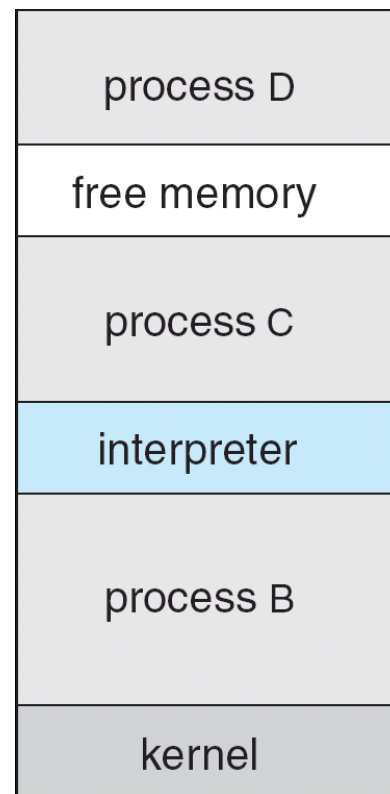
(a)



(b)

■ MS-DOS

(a) At system startup (b) running a program



■ Unix

Shell (interpreter) as one of user processes





Bourne Shell Command Interpreter

```

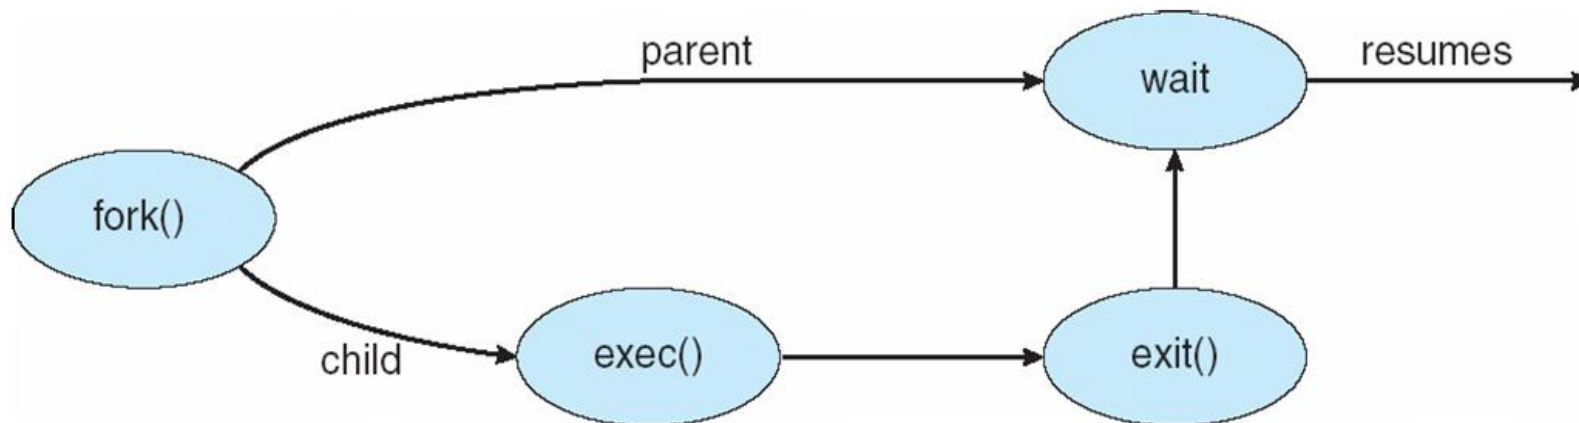
Terminal
File Edit View Terminal Tabs Help
fd0      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
sd0      0.0    0.2    0.0    0.2  0.0  0.0    0.4  0  0
sd1      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
          extended device statistics
device   r/s    w/s    kr/s    kw/s wait actv  svc_t  %w  %b
fd0      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
sd0      0.6    0.0   38.4    0.0  0.0  0.0    8.2  0  0
sd1      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
(root@pbg-nv64-vn)-(11/pts)-(00:53 15-Jun-2007)-(global)
- (/var/tmp/system-contents/scripts)# swap -sh
total: 1.1G allocated + 190M reserved = 1.3G used, 1.6G available
(root@pbg-nv64-vn)-(12/pts)-(00:53 15-Jun-2007)-(global)
- (/var/tmp/system-contents/scripts)# uptime
12:53am up 9 min(s), 3 users, load average: 33.29, 67.68, 36.81
(root@pbg-nv64-vn)-(13/pts)-(00:53 15-Jun-2007)-(global)
- (/var/tmp/system-contents/scripts)# w
4:07pm up 17 day(s), 15:24, 3 users, load average: 0.09, 0.11, 8.66
User      tty          login@ idle   JCPU   PCPU   what
root      console      15Jun07 18days 1      /usr/bin/ssh-agent -- /usr/bi
n/d
root      pts/3        15Jun07 18      4      w
root      pts/4        15Jun07 18days 4      w
(root@pbg-nv64-vn)-(14/pts)-(16:07 02-Jul-2007)-(global)
- (/var/tmp/system-contents/scripts)#

```





Process Creation in a Nutshell





Linux Shell

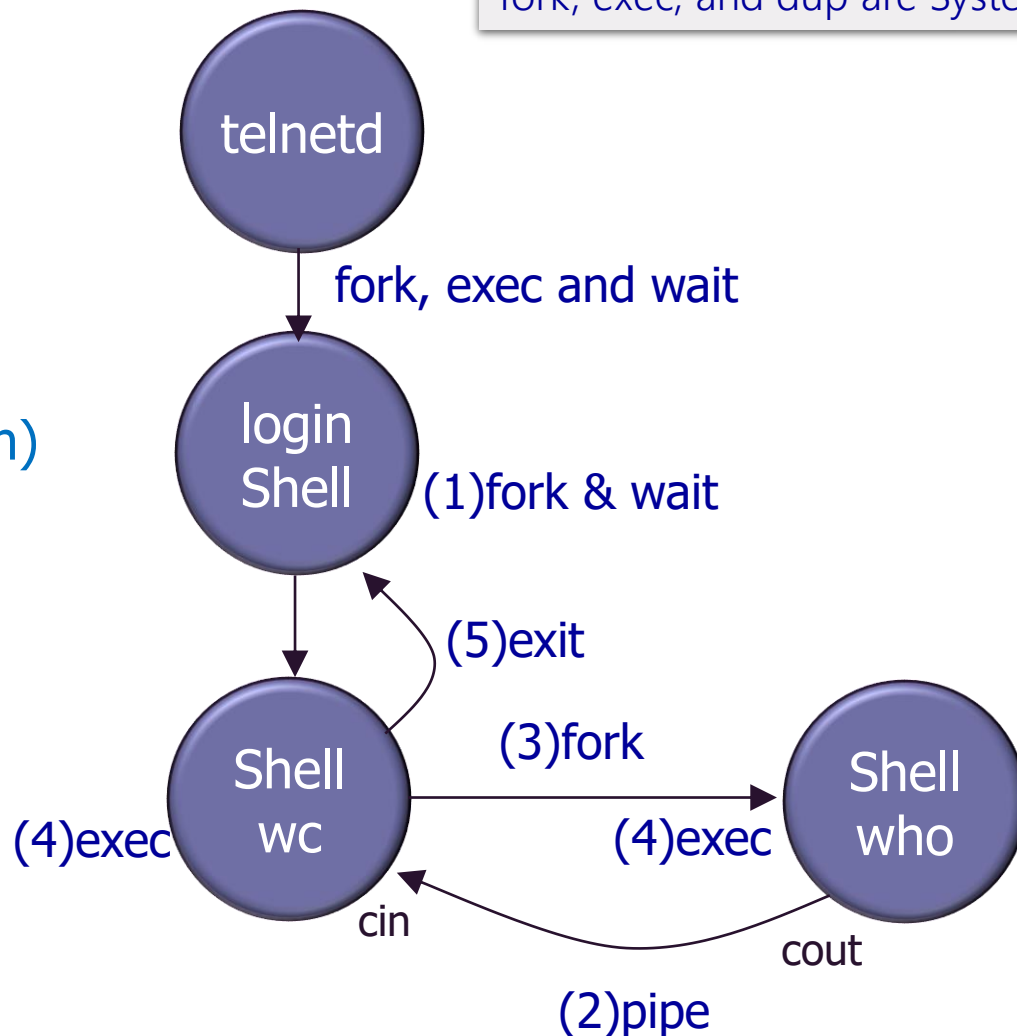
fork, exec, and dup are System calls



goodall login: **chinchia**

goodall[1]% **(you type sth)**

goodall[1]% **who | wc -l**





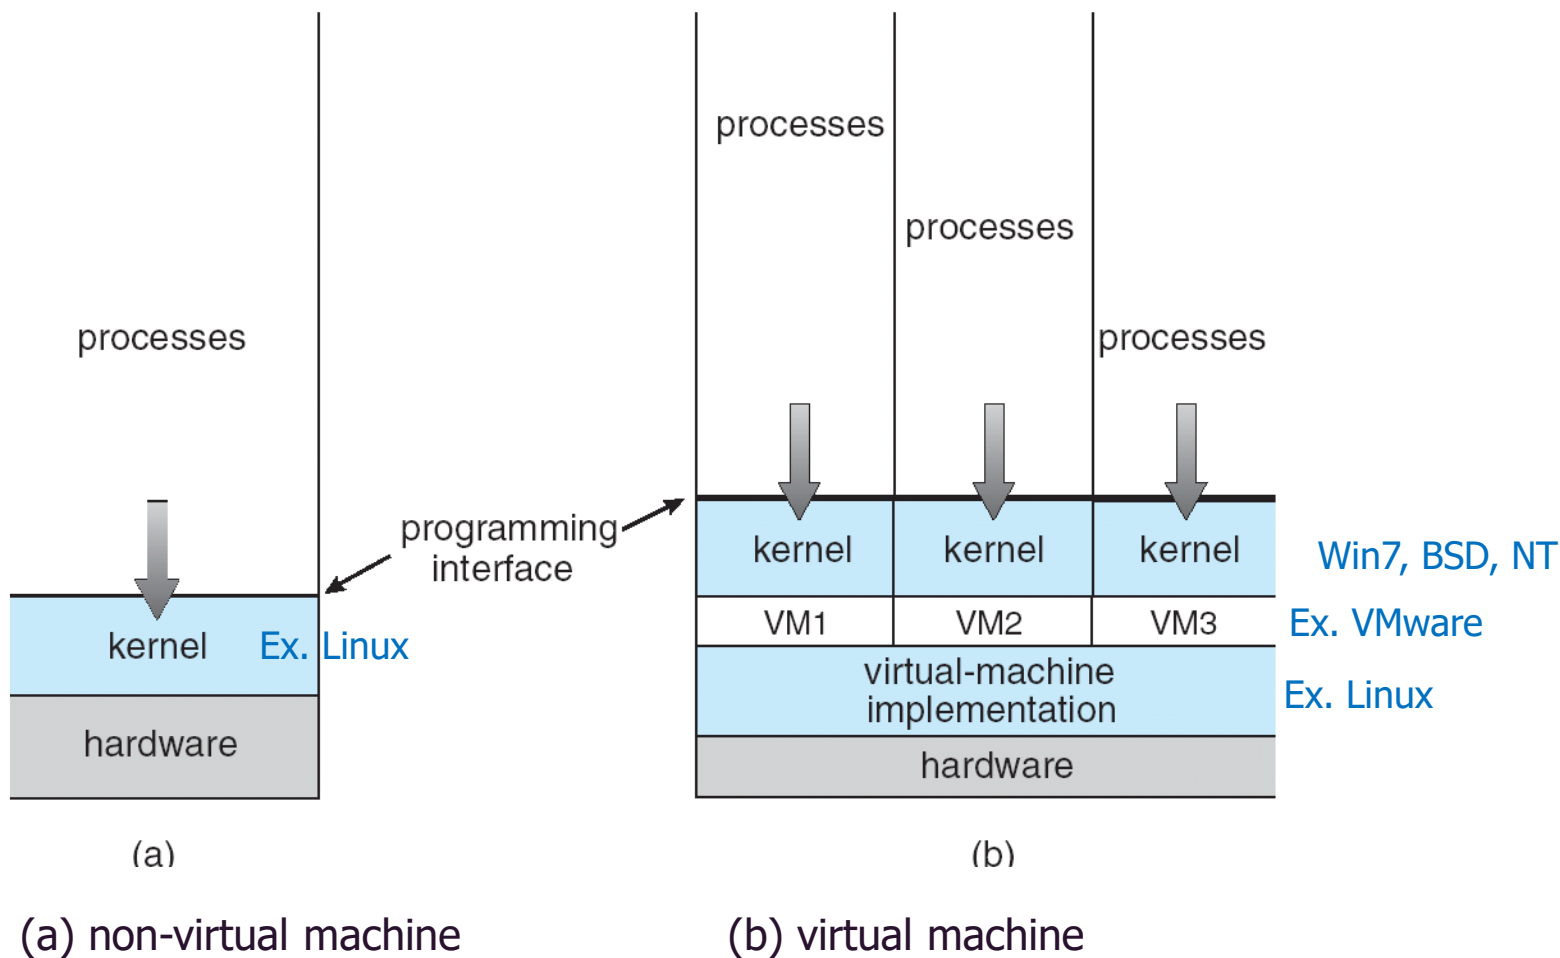
Virtual Machines

- A **virtual machine** takes the **layered approach** to its logical conclusion. It treats hardware and the operating system kernel as though they were all hardware.
- A virtual machine provides an interface *identical* to the underlying bare hardware.
- The operating system **host** creates the illusion that a process has its own processor and (virtual memory).
- Each **guest** is provided with a (virtual) copy of underlying computer.



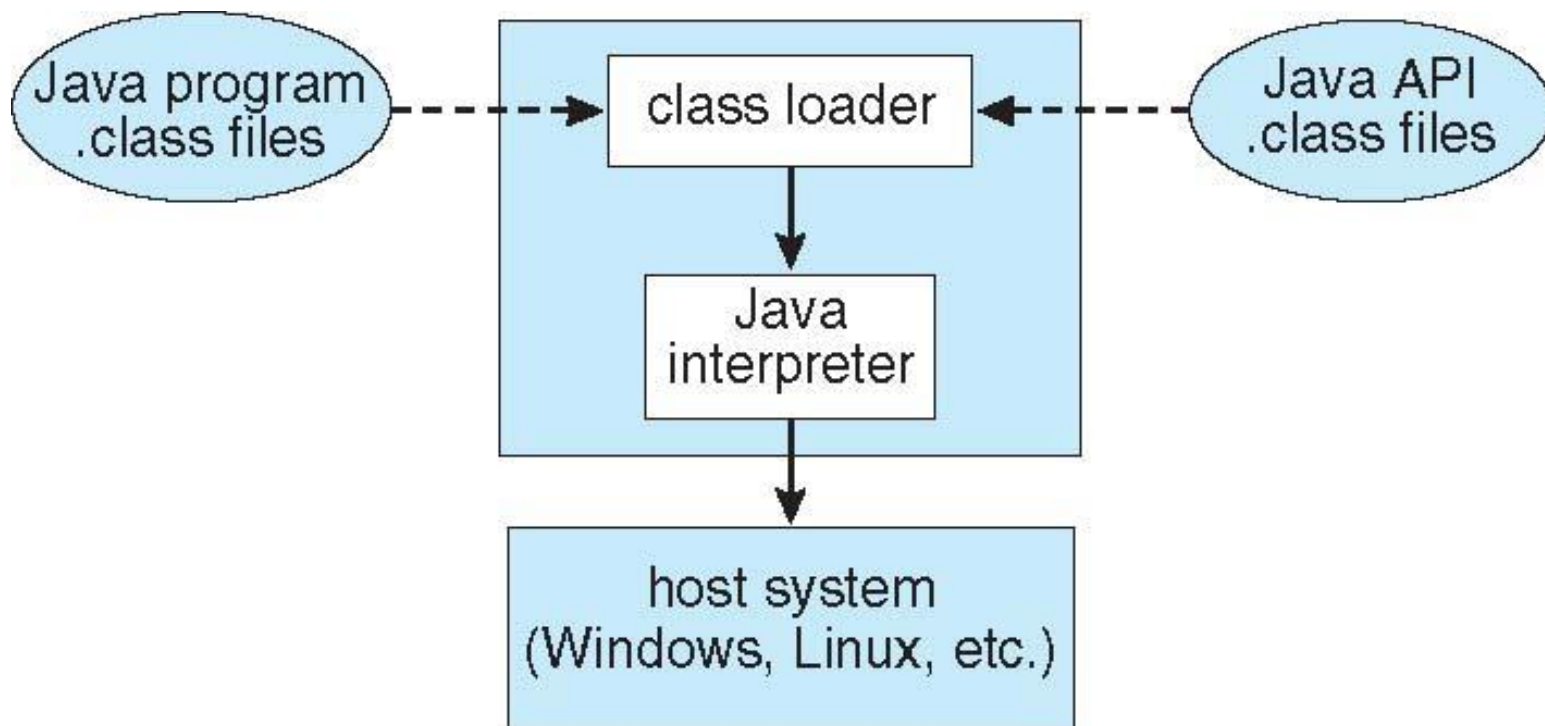


Virtualization





Java Virtual Machine





Discussions 1: TBD

1. In what particular situation have your program received a *segmentation fault*?
2. To read data from a **file**, why do we need to call open and close the file? In other words, why doesn't OS allow read(filename, data, size)?
3. If your C++ program terminates upon an **exception**, it may not print out a **cout** statement that must have been executed before the exception. Why?





System Boot

- An operating system must be made available to hardware so hardware can start it.
 - Small piece of code – **bootstrap loader**, locates the kernel, loads it into memory, and starts it.
 - Sometimes two-step process where **boot block** at fixed location loads bootstrap loader.
 - When power initialized on system, execution starts at a fixed memory location.
 - Firmware is used to hold initial boot code.

